Manya Jain

Tempe, AZ | manyajainrkm@gmail.com | +1 480-637-1065 | linkedin.com/in/manya-jainn | github.com/ManyaJainrkm

Education

Arizona State University Tempe, AZ
Master of Science in Computer Science Expected May 2027
GPA: 3.50/4.00

Manipal University Jaipur India
Bachelor of Technology in Computer Science Engineering Aug 2021 - May 2025

Research: Co-Author, *Mathematical Models of AI and Its Legal and Ethical Boundaries* (2023)

Skills

Languages: Python, SQL, Java, C, C++, Bash
AI/ML & GenAI: pandas, scikit-learn, XGBoost, LangChain, RAG, Snowflake Cortex, LLM evaluation
Data & Analytics: statistical modeling, fairness/bias analysis, data validation, Tableau
Systems & Tools: Linux, Docker, Kubernetes, Jenkins, CI/CD, REST APIs, Git
Databases: PostgreSQL, MySQL, Neo4j, Redis

Experience

CIPS-AI Lab, Arizona State University Tempe, AZ
Student Researcher, Applied AI Systems Feb 2026 - Present

- Identified and validated bigram and trigram feature extraction for legal NLP classification, enabling the model to recognize multi-word domain-specific signals like 'consent judgment' as semantic units
- Built Python ETL pipelines to clean and structure large volumes of unstructured documents for downstream ML classification, enforcing reproducibility through structured logging across 20+ evaluation runs
- Prototyped Neo4j knowledge graphs linking cases to support Cypher-based querying and retrieval-augmented workflows

Safexpress Private Limited Gurugram, India
Software Engineer Intern, GenAI May 2026 - Aug 2026

- Evaluated a Snowflake-based Gen AI business intelligence assistant, systematically testing outputs against ground-truth sales data across India's largest B2B logistics network to surface accuracy failures
- Found that mandatory data-integrity filters excluding cancelled/void transactions were applied inconsistently, producing confident but incorrect figures; restructured rule enforcement to prevent invalid records from corrupting reported numbers
- Diagnosed geographic disambiguation errors where queries silently resolved to the wrong data scope, proposing fixes to ensure outputs reflected the intended boundaries

National Institute of Urban Affairs, Govt. of India New Delhi, India
Software Engineering Intern, DevOps & Platform Systems Aug 2024 - May 2025

- Built and maintained Jenkins CI/CD pipelines across 12+ containerized services, improving release reliability by 30% through build validation, configuration checks, and deployment gating
- Implemented Prometheus and Grafana monitoring stack to enable proactive incident detection across deployments
- Reduced incident recovery time by 25% by tracing Linux logs, container states, SQL queries, and OAuth failures to isolate root causes and restore service within SLA

Ministry of Housing and Urban Affairs New Delhi, India
Data Analytics Intern Dec 2023 - Feb 2024

- Built Tableau dashboards and visualizations analyzing PAN-India urban development data spanning multiple years, including government-facing reporting deliverables and public-facing visualizations
- Applied advanced Excel techniques including pivot tables, VLOOKUP/XLOOKUP, and conditional formatting to clean, structure, and analyze large-scale national datasets.

Siemens Healthineers India New Delhi, India
Summer Intern, Digital Workflows May 2023 - Jul 2023

- Analyzed inspection activity across 100+ staff using SafetyCulture, tracked year-over-year trends to identify engagement patterns and inform outreach strategy
- Designed a segmented employee outreach workflow based on activity thresholds

Projects

Fairness-Aware Income Classification System Jan 2026 - May 2026
Python, scikit-learn, pandas, XGBoost, Git

- Built an end-to-end ML pipeline on 32,000+ census records achieving ROC-AUC $>$ 0.85 with Logistic Regression and XGBoost, implementing demographic parity and equalized odds checks to quantify model bias and fairness trade-offs

HireMate: Graph-Augmented Job Matching System Aug 2025 - Dec 2025
Python, Neo4j, SQL, LangChain, Embedding Retrieval

- Built a RAG-based backend combining Neo4j graph traversal with embedding-based ranking across 5,000+ job profiles using LangChain-orchestrated retrieval pipelines, reducing query latency by 25% via indexing and Cypher optimization